

DATASCI W261: Machine Learning at Scale

David Rose
david.rose@berkeley.edu
W261-1
Week 01
2015.08.31

HW1.0.0

Everyone seems to have their own definition of big data, the 3 Vs, etc. Here's another: big data is information of sufficient size and complexity to require a new and different set of tools and techniques to effectively make use of it as compared to traditional data processing. A corollary of this definition is that in the near future what is considered big data will no longer be, since the tools and techniques will have become the new normal.

The human genome data is an example of big data. Genomic information will play an increasingly large role in healthcare and population health in the future. In terms of size the 1000 Genomes Project contains more than 200 TB of data, for example.

HW1.0.1

Let n be the number of polynomial regression models to be considered.

Let m_n be the n^{th} model with polynomial degree n .

Let j be the number of records in dataset T .

Let k be the number of desired subsets of T to be used for training and testing.

Divide T into k subsets containing (j/k) records in each, such that T_k is the j^{th} subset of T .

For each model m_p do

 For each data subset T_q do

 Divide T_q into two non-overlapping subsets of equal size, T_q^{train} and T_q^{test} , to be used for training and testing model m_p

 Train model m_p on data T_q^{train}

For each tuple $[x,y]$ in $T_{q^{test}}$

Calculate $m_p(x)$; store $m_p(x)$ and y for subsequent calculations

For each model calculate estimated average bias, variance and prediction error:

$bias^2$ = average of the average squared difference between $m_p(x)$ and y for each $T_{q^{test}}$

variance = average difference between $m_p(x)$ and the average value of $m_p(x)$ for all datasets $T_{q^{test}}$

prediction error = $bias^2$ + variance plus a constant representing noise for each $T_{q^{test}}$. The constant is ignored since it is, after all, constant.

The best model is selected by determining the model with the minimal prediction error.

HW1.1

```
In [71]: 1 #HW1.1. Read through the provided control script (pNaiveBayes.sh)
         2 !printf 'done'
```

done

Mapper script. The same mapper is used for each subsequent exercise. Of note:

- both email subject and email body are considered together
- in addition to emitting word counts the mapper also emits email classification counts
 - this is not consistent with functional programming, but here it's okay

```

In [58]: 1 %%writefile mapper.py
2          #!/usr/bin/python
3          ''' mapper reads name of file containing chunk of email records and
4             words of interest
5
6             mapper emits counts of words, and counts of email classification
7          '''
8          from __future__ import print_function
9          import re
10         import string
11         import sys
12         filename = sys.argv[1]
13         findwords = sys.argv[2].split()
14         # regular expression to remove all punctuation
15         punctuation = re.compile('[%s]' % re.escape(string.punctuation))
16         with open (filename, "r") as myfile:
17             for line in myfile:
18                 # split line into three tokens: id, classification, email
19                 # both the email subject and the email body are included in
20                 tokens = line.split('\t', 2)
21                 isspam = tokens[1]
22                 # emit count of email classification, using magic word '__
23                 if isspam == '0': # ham
24                     print('__CLASS__', 1, 0)
25                 else: # spam
26                     print('__CLASS__', 0, 1)
27                 # convert text to lower case
28                 text = tokens[len(tokens) - 1].lower()
29                 # remove punctuation from text
30                 text = punctuation.sub('', text)
31                 # split into individual words
32                 words = re.findall(r"[\w']+", text)
33                 for word in words:
34                     # only report on word if it is in the word list parameter
35                     # or report on all words if parameter equals '*'
36                     if word in findwords or sys.argv[2] == '*':
37                         # emit the word and the classification count
38                         if isspam == '0': # ham
39                             print(word, 1, 0)
40                         else: # spam
41                             print(word, 0, 1)
42

```

Overwriting mapper.py

Reducer script for HW1.2. Aggregates results from mapper, emits summed counts of all words processed by mapper.

```
In [59]: 1 %%writefile reducer.py
2 #!/usr/bin/python
3 ''' reducer is provided list of temporary files containing mapper
4
5     reducer reads each file and aggregates counts of words, then e
6 '''
7 from __future__ import print_function
8 import sys
9 filelist = sys.argv
10 words = {}
11 while len(filelist) > 1: # do not use sys.argv[0]
12     with open(filelist.pop(), 'r') as cfile:
13         for line in cfile:
14             tokens = line.split()
15             word = tokens[0]
16             if word not in words.keys():
17                 words[word] = 0
18             # mapper produces counts based on email classification
19             # this reducer is only interested in total counts
20             words[word] += int(tokens[1]) + int(tokens[2])
21 # emit results
22 for word in sorted(words.keys()):
23     # ignore counts for email classification
24     if word != '__CLASS__':
25         print('\t'.join([word, str(words[word])]), file=sys.stderr)
26         print('\t'.join([word, str(words[word])]))
27
```

Overwriting reducer.py

HW1.2

```
In [60]: 1 # HW1.2. Provide a mapper/reducer pair that, when executed by pNaive
2 # will determine the number of occurrences of a single, user-speci
3 !chmod +x *.py
4 !./pNaiveBayes.sh 4 "assistance"
```

```
assistance      10
```

Reducer script for HW1.3-5. Aggregates results from mapper to generate vocabulary and email classification counts.

Test the results against the same data set, classifying email records and comparing the results to known classifications.

NOTE: executing the next cell will overwrite the reducer script created in the earlier cell

```
In [61]: 1 %%writefile reducer.py
2 #!/usr/bin/python
3 ''' reducer is provided a list of temporary files containing mapper
4
```

```

5     reducer reads each file and aggregates counts of words and ema
6
7     reducer then applies a Naive Bayes classifier against the same
8     to build the training parameters, classifying email records and
9     results to known classifications.
10 '''
11 from __future__ import print_function
12 import math
13 import re
14 import string
15 import sys
16 # store statistics on the original list of words of interest
17 keywords = {}
18 # counts of each email classification
19 hamcount = 0
20 spamcount = 0
21 # counts of words in each email classification
22 spamwordcount = 0
23 hamwordcount = 0
24
25 filelist = sys.argv
26 while len(filelist) > 1:
27     with open(filelist.pop(), 'r') as cfile:
28         for line in cfile:
29             tokens = line.split()
30             word = tokens[0]
31             # special case for count of email classification
32             if word == '__CLASS__':
33                 hamcount += int(tokens[1])
34                 spamcount += int(tokens[2])
35             # regular case of count of word
36             else:
37                 if word not in keywords.keys():
38                     keywords[word] = [0, 0]
39                     keywords[word][0] += int(tokens[1])
40                     keywords[word][1] += int(tokens[2])
41                 hamwordcount += int(tokens[1])
42                 spamwordcount += int(tokens[2])
43 # total number of unique words
44 vocabcount = len(keywords)
45 # total number of email records
46 doccount = spamcount + hamcount
47
48 # counters for determining error rate
49 correct = 0
50 incorrect = 0
51
52 # regular expression for removing punctuation
53 punctuation = re.compile('[%s]' % re.escape(string.punctuation))
54 with open('enronemail_1h.txt', 'r') as cfile:
55     for line in cfile:
56         # words to be used in Naive Bayes classification
57         nbwords = {}
58         tokens = line.split('\t', 2)

```

```

59     eid = tokens[0]
60     issspam = tokens[1]
61     # build bag of words for email record
62     text = tokens[len(tokens) - 1].lower()
63     text = punctuation.sub('', text)
64     docwords = re.findall(r"\w+", text)
65     for word in docwords:
66         if word in keywords.keys():
67             if word not in nbwords:
68                 nbwords[word] = 1
69             else:
70                 nbwords[word] += 1
71
72     # calculate the probability of the email record being spam
73     # natural log conversion is used to avoid floating point
74
75     # start with the prior probability of a spam record
76     logspam = math.log(spamcount / float(doccount))
77     for word in nbwords:
78         # add the probability of the word being present in the
79         # multiplied by the number of times the word appears
80         logspam += (nbwords[word] *
81                     (math.log(keywords[word][1] + 1 / float(spamwordcount)
82
83     # start with the prior probability of a ham record
84     logpham = math.log(hamcount / float(doccount))
85     for word in nbwords:
86         # add the probability of the word being present in the
87         # multiplied by the number of times the word appears
88         logpham += (nbwords[word] * (math.log(keywords[word][1] + 1 / float(hamwordcount)
89
90     # determine the classification, based on comparison of log
91     nbclass = '0'
92     if logspam > logpham:
93         nbclass = '1'
94
95     # add some statistics
96     if issspam == nbclass:
97         correct += 1
98     else:
99         incorrect += 1
100
101     # emit the results
102     #print('\t'.join([eid, issspam, nbclass, str(issspam == nbclass)]))
103     print(''\t'.join([eid, issspam, nbclass]))
104 # print some statistics
105 print('correct: {}, incorrect: {}, training error: {}'.format(correct,
106     str(float(incorrect) / (correct + incorrect))), file=sys.stderr)
107
108

```

Overwriting reducer.py

```
In [62]: 1 !chmod +x *.py
```

HW1.3

```
In [63]: 1 # HW1.3. Provide a mapper/reducer pair that, when executed by pNai
2 # will classify the email messages by a single, user-specified wor
3 # using the Naive Bayes Formulation.
4 !./pNaiveBayes.sh 4 "assistance"
```

correct: 60, incorrect: 40, training error: 0.4

HW1.4

```
In [64]: 1 # HW1.4. Provide a mapper/reducer pair that, when executed by pNai
2 # will classify the email messages by a list of one or more user-s
3 !./pNaiveBayes.sh 4 "assistance valium enlargementWithATypo"
```

correct: 63, incorrect: 37, training error: 0.37

HW1.5

```
In [65]: 1 # HW1.5. Provide a mapper/reducer pair that, when executed by pNai
2 # will classify the email messages by all words present.
3 !./pNaiveBayes.sh 4 ""
```

correct: 100, incorrect: 0, training error: 0.0

The benchmark script compares performance (in terms of error rates, not execution time) of the SciKit-Learn implementations of the Multinomial Naive Bayes algorithm and the Bernoulli Naive Bayes algorithm.

In [66]:

```
1 %%writefile benchmark.py
2 #!/Users/david/anaconda/bin/python
3 from __future__ import print_function
4 import re
5 import string
6 from sklearn.naive_bayes import BernoulliNB
7 from sklearn.naive_bayes import MultinomialNB
8 from sklearn.feature_extraction.text import CountVectorizer
9 import sys
10 records = []
11 labels = []
12 # regular expression for removing punctuation
13 punctuation = re.compile('%s' % re.escape(string.punctuation))
14
15 # read the input data and create separate lists for content and cl
16 with open('enronemail_1h.txt', 'r') as cfile:
17     for line in cfile:
18         tokens = line.split('\t', 2)
19         eid = tokens[0]
20         label = tokens[1]
21         # prepare text
22         text = tokens[len(tokens) - 1].lower()
23         text = punctuation.sub('', text)
24         records.append(text) # content
25         labels.append(label) # classification
26 # prepare the features, using the SciKit-Learn CountVectorizer
27 data = CountVectorizer().fit_transform(records)
28
29 # train and test using the Multinomial Naive Bayes implementation
30 clf = MultinomialNB()
31 clf.fit(data, labels)
32 results = clf.predict(data)
33 # measure and report training error
34 incorrect = 0
35 for a,b in zip(labels, results):
36     incorrect += not a == b
37 print('Multinomial NB Training Error: ', str(float(incorrect) / len(labels)))
38
39 # train and test using the Multinomial Naive Bayes implementation
40 clf = BernoulliNB()
41 clf.fit(data, labels)
42 results = clf.predict(data)
43 # measure and report training error
44 incorrect = 0
45 for a,b in zip(labels, results):
46     incorrect += not a == b
47 print('Bernoulli NB Training Error: ', str(float(incorrect) / len(labels)))
48
49
50
```

Overwriting benchmark.py

HW1.6

```
In [73]: 1 # HW1.6 Benchmark your code with the Python SciKit-Learn implement
        2 !chmod +x *py
        3 !./benchmark.py
        4 !printf 'HW1.5 Training Error: ' && ./pNaiveBayes.sh 4 ""
```

Multinomial NB Training Error: 0.0

Bernoulli NB Training Error: 0.21

HW1.5 Training Error: correct: 100, incorrect: 0, training error: 0.0

Results:

Model Type	Training Error
Multinomial NB	0.0
Bernoulli NB	0.21
HW1.5	0.0

Discussion: There are no differences in the results between the SciKit-Learn Multinomial Naive Bayes implementation and the HW1.5 implementation. Since both are training and testing over the same data set it is not surprising that both achieve a training error rate of 0.0.

As seen in the table above, the SciKit-Learn Bernoulli Naive Bayes implementation did not perform as well as the Multinomial Naive Bayes implementation. This can be ascribed to the fact that the Bernoulli approach uses a dichotomous value for the presence or absence of a term in an email record, whereas the Multinomial approach takes into consideration the number of times a term occurs, yielding a more accurate representation.